

Homework 1 - Practical component

- This homework must be done and submitted to Gradescope individually. You are welcome to discuss with other students but the solution and code you submit must be your own. Note that we will use Gradescope's plagiarism detection feature. All suspected cases of plagiarism will be recorded and shared with university officials for further handling.
- The practical part should be coded in python (you can use the numpy and matplotlib libraries) and all code will be submitted as a python file to Gradescope. To enable automated code grading you should work off of the template file given in this homework folder. Do not modify the name of the file or any of the function signatures of the template file or the code grading will not work for you. You may, of course, add new functions and any regular python imports.
- Any graphing, charts, or practical report parts should be submitted in a pdf to Gradescope. For the report it is recommended to use a Jupyter notebook, writing math with MathJax and export as pdf. You may alternatively write your report in L^AT_EX; L^AT_EX; Word. In any case, you should export your report to a pdf file that you will submit. You are of course encouraged to draw inspiration from what was done in lab sessions.

You should work off of the template file `solution.py` in the project and fill in the basic numpy functions using numpy and python methods

Parzen with soft windows (kernels)

In this Homework we will use Iris Data Set as a toy dataset. It contains 150 samples (one for each row), each with 4 features (the 4 first columns) and

one label in $\{1, 2, 3\}$ (the 5th column). It is recommended you test your code by importing it like this :

```
import numpy as np
iris = np.genfromtxt('iris.txt')
```

When the answer template in `solution.py` has "iris" as an argument, you may assume that this argument is the dataset in numpy format. Your function should use this argument to perform computations, not a version of the dataset that you loaded by yourself.

1. [4 points]

Question. Write functions that take that dataset as input and return the following statistics:

- (a) `Q1.feature_means` : An array containing the empirical means of each feature, from all examples present in the dataset. Make sure to maintain the original order of features.

e.g. : `Q1.feature_means(iris) = [$\mu_1, \mu_2, \mu_3, \mu_4$]`

- (b) `Q1.covariance_matrix` : A 4×4 matrix that represents the empirical covariance matrix of features on the whole dataset.

- (c) `Q1.feature_means_class_1` : An array containing the empirical means of each feature, but only from examples in class 1. The possible classes in the iris dataset are 1, 2, and 3.

e.g. : `Q1.feature_means_class_1(iris) = [$\mu_1, \mu_2, \mu_3, \mu_4$]`

- (d) `Q1.covariance_matrix_class_1` : A 4×4 matrix that represents the empirical covariance matrix of features, but only from examples in class 1.

2. [1 points]

Question. Implement Parzen with hard window parameter h . Use the standard Euclidean distance on the original features of the dataset.

Recall that the Parzen Window for Multi-class classification $Y_i \in \{1, \dots, m\}$ is defined as follows:

$$f(x) = \operatorname{argmax} \left(\frac{1}{\sum_{i=1}^n K(X_i, x)} \sum_{i=1}^n K(X_i, x) \operatorname{onehot}_m(Y_i) \right)$$

and the kernel for a hard window with parameter h is the following:

$$K(X_i, x) = \mathbf{1}_{\{d(X_i, x) < h\}}$$

Your answer should have the following behavior :

f = HardParzen(h) initiates the algorithm with parameter h

f.train(X, Y) trains the algorithm, where X is a $n \times m$ matrix of n training samples with m features, and Y is an array containing the n labels. The labels are denoted by integers, but the number of classes in Y can vary.

f.compute_predictions(X_test) computes the predicted labels and return them as an array of same size as X_test . X_test is a $k \times m$ matrix of k test samples with same number of features as X . This function is called only after training on (X, Y) . **If a test sample x has no neighbor within window h , the algorithm should choose a label at random by using `draw_rand_label(x, label_list)`, a function that is provided in the `solution.py` file, where `label_list` is the list of different classes present in Y , and x is the array of features of the corresponding point.**

3. [5 points]

Question. Implement Parzen with a soft window. We will use a radial basis function (RBF) kernel (also known as *Gaussian* kernel) with parameter σ .

Recall that the RBF kernel is defined as follows:

$$K(X_i, x) = \mathcal{N}_{X_i, \sigma^2}(x) = \frac{1}{(2\pi)^{\frac{D}{2}} \sigma^D} e^{-\frac{1}{2} \frac{d(X_i, x)^2}{\sigma^2}}$$

where D is the number of features of the dataset and d is the Euclidean distance.

The structure of your solution should be the same as in the previous question, but you will never need to draw a label at random with `draw_rand_label(x, label_list)`. The class name is **SoftRBF-Parzen**.

4. [5 points]

Question. Implement a function `split_dataset` that splits the iris dataset as follows:

- A training set consisting of the samples of the dataset with indices which have a remainder of either 0 or 1, or 2 when divided by 5

- A validation set consisting of the samples of the dataset with indices which have a remainder of 3 when divided by 5.
- A test set consisting of the samples of the dataset with indices which have a remainder of 4 when divided by 5.

For instance the element of index 14 (in the original dataset) should be part of the test set because the remainder of 14 divided by 5 is 4. Do not use random splitting for this exercise (even though it is generally a very good idea). The function should take as input the dataset and return the three sets as a tuple (train, validation, test), where each element of the tuple is a matrix with 5 columns (the 4 features and the labels, kept in the same order).

5. [10 points]

Question. Implement two functions **ErrorRate.hard_parzen** and **ErrorRate.soft_parzen** that compute the error rate (i.e. the proportion of missclassifications) of the HardParzen and SoftRBFParzen algorithms. The expected behavior is as follows :

test_error = ErrorRate(x_train, y_train, x_val, y_val) initiates the class and stores the training and validation sets, where **x_train** and **x_val** are matrices with 4 feature columns, and **y_train** and **y_val** are arrays containing the labels.

test_error.hard_parzen(h) takes as input the window parameter **h** and returns as a float the error rate on **x_val** and **y_val** of the Hard-Parzen algorithm that has been trained on **x_train** and **y_train**.

test_error.soft_parzen(σ) works just like with Hard Parzen, but with the SoftRBFParzen algorithm.

Then, include in your report a single plot with two lines:

- (a) Hard Parzen window's classification error on the validation set of iris, when trained on the training set (see question 4) for the following values of **h**:

$$h \in \{0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 1.0, 3.0, 10.0, 20.0\}$$

- (b) RBF Parzen's classification error on the validation set of iris, when trained on the training set (see question 4) for the following values of σ :

$$\sigma \in \{0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 1.0, 3.0, 10.0, 20.0\}$$

The common x-axis will represent either h or σ . Always label your axes and lines in the plot!

Give a detailed discussion of your observations.

6. [5 points]

Question. Implement a function `get_test_errors` that uses the evaluated validation errors from the previous question to select h^* and σ^* , then computes the error rates on the test set. The value h^* is the one (among the proposed set in question 5) that results in the smallest validation error for Parzen with hard window, and σ^* is the parameter (among the proposed set in question 5) that results in the smallest validation error for Parzen with RBF.

The function should take as input the dataset and split it using question 4. The expected output is an array of size 2, the first value being the error rate on the test set of Hard Parzen with parameter h^* (trained on the training set), and the second value being the error rate on the test set of Soft RBF Parzen with parameter σ^* (trained on the training set).

7. [5 points]

Question. Include in your report a discussion on the running time complexity of these two methods. How does it vary for each method when the hyperparameter h or σ changes? Why ?

8. [5 points]

Question. Implement a random projection (Gaussian sketch) map to be used on the input data:

Your function `random_projection` should accept as input a feature matrix X of dimension $n \times 4$, as well as a 4×2 matrix A encoding our projection.

Define $p : x \mapsto \frac{1}{\sqrt{2}}xA$ and use this random projection map to reduce the dimension of the inputs (feature vectors of the dataset) from 4 to 2.

Your function should return the output of the map p when applied to X , in the form of a $n \times 2$ matrix.

e.g. $random_projection(X_{n,4}, A_{4,2}) = X_{n,2}^{proj}$

9. [10 points]

Question. Similar to Question 5, compute the validation errors of

Hard Parzen classifiers trained on 500 random projections of the training set, for

$$h \in \{0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 1.0, 3.0, 10.0, 20.0\}$$

The validation errors should be computed on the projected validation set, using the same matrix A . To obtain random projections, you may draw A as 8 independent variables drawn uniformly from a gaussian distribution of mean 0 and variance 1.

You can for example store these validation errors in a 500×10 matrix, with a row for each random projection and a column for each value of h .

Do the same thing for RBF Parzen classifiers, for

$$\sigma \in \{0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 1.0, 3.0, 10.0, 20.0\}$$

Plot and include in your report in the same graph the average values of the validation errors (over all random projections) for each value of h and σ , along with error bars of length equal to $0.2 \times$ the standard deviations.

How do your results compare to the previous ones?